

# Simple Linear Regression

## What I'm Doing

I implemented a simple linear regression with scikit-learn, using engine size to predict **CO<sub>2</sub> emissions** from Canadian vehicle fuel data. I split the dataset 80/20 for training and testing, then evaluated the model with MAE, MSE, and R<sup>2</sup>.

---

## 1. Setting Up

I start by importing the essential libraries for data manipulation, numpy, and visualization.

```
%pip install matplotlib pandas numpy scikit-learn
import matplotlib.pyplot as plt
import pandas as pd
import pylab as pl
import numpy as np
%matplotlib inline
```

download dataset from IBM Object Storage using `!curl -o` (Linux: `!wget -O`)

```
!curl -o FuelConsumption.csv https://cf-courses-data.s3.us.cloud-object-
storage.appdomain.cloud/IBMDeveloperSkillsNetwork-ML0101EN-SkillsNetwork/labs/
Module%202/data/FuelConsumptionCo2.csv
```

| % Total | % Received | % Xferd | Average Speed | Time  | Time     | Time     | Current                          |
|---------|------------|---------|---------------|-------|----------|----------|----------------------------------|
|         |            |         | Dload Upload  | Total | Spent    | Left     | Speed                            |
| 0       | 0          | 0       | 0             | 0     | --:--:-- | --:--:-- | 0                                |
| 0       | 0          | 0       | 0             | 0     | --:--:-- | --:--:-- | 0                                |
| 0       | 0          | 0       | 0             | 0     | --:--:-- | 0:00:01  | 0                                |
| 100     | 72629      | 100     | 72629         | 0     | 0        | 31942    | 0 0:00:02 0:00:02 --:--:-- 31981 |

---

## 2. Getting to Know the Dataset

Then I load and inspect the first few rows.

```
df = pd.read_csv("FuelConsumption.csv")
df.head()
```

|   | MODELYEAR | MAKE  | MODEL         | VEHICLECLASS | ENGINESIZE | CYLINDERS | TRANSMISSION | FUEL' |
|---|-----------|-------|---------------|--------------|------------|-----------|--------------|-------|
| 0 | 2014      | ACURA | ILX           | COMPACT      | 2.0        | 4         |              | AS5   |
| 1 | 2014      | ACURA | ILX           | COMPACT      | 2.4        | 4         |              | M6    |
| 2 | 2014      | ACURA | ILX<br>HYBRID | COMPACT      | 1.5        | 4         |              | AV7   |
| 3 | 2014      | ACURA | MDX<br>4WD    | SUV - SMALL  | 3.5        | 6         |              | AS6   |
| 4 | 2014      | ACURA | RDX<br>AWD    | SUV - SMALL  | 3.5        | 6         |              | AS6   |

**Result:** The dataset contains model year, make, model, vehicle class, engine specs, fuel consumption ratings, and the target - CO<sub>2</sub> emissions in g/km. [Dataset source](#)

A quick statistical summary of each numerical column.

```
df.describe()
```

|       | MODELYEAR | ENGINESIZE  | CYLINDERS   | FUELCONSUMPTION_CITY | FUELCONSUMPTION_HWY |
|-------|-----------|-------------|-------------|----------------------|---------------------|
| count | 1067.0    | 1067.000000 | 1067.000000 | 1067.000000          | 1067.000000         |
| mean  | 2014.0    | 3.346298    | 5.794752    | 13.296532            | 9.474602            |
| std   | 0.0       | 1.415895    | 1.797447    | 4.101253             | 2.794510            |
| min   | 2014.0    | 1.000000    | 3.000000    | 4.600000             | 4.900000            |
| 25%   | 2014.0    | 2.000000    | 4.000000    | 10.250000            | 7.500000            |
| 50%   | 2014.0    | 3.400000    | 6.000000    | 12.600000            | 8.800000            |
| 75%   | 2014.0    | 4.300000    | 8.000000    | 15.550000            | 10.850000           |
| max   | 2014.0    | 8.400000    | 12.000000   | 30.200000            | 20.500000           |

### 3. Exploring the Data


I narrow my focus to the four most relevant features for this regression task.

```
cdf = df[['ENGINE SIZE', 'CYLINDERS', 'FUELCONSUMPTION_COMB', 'CO2EMISSIONS']]
cdf.head(9)
```

|   | ENGINE SIZE | CYLINDERS | FUEL CONSUMPTION_COMB | CO2 EMISSIONS |
|---|-------------|-----------|-----------------------|---------------|
| 0 | 2.0         | 4         | 8.5                   | 196           |
| 1 | 2.4         | 4         | 9.6                   | 221           |
| 2 | 1.5         | 4         | 5.9                   | 136           |
| 3 | 3.5         | 6         | 11.1                  | 255           |
| 4 | 3.5         | 6         | 10.6                  | 244           |
| 5 | 3.5         | 6         | 10.0                  | 230           |
| 6 | 3.5         | 6         | 10.1                  | 232           |
| 7 | 3.7         | 6         | 11.1                  | 255           |
| 8 | 3.7         | 6         | 11.6                  | 267           |

Histograms give a quick sense of how each feature is distributed.


```
viz = cdf[['CYLINDERS', 'ENGINE SIZE', 'CO2 EMISSIONS', 'FUEL CONSUMPTION_COMB']]
viz.hist()
plt.show()
```



**Result:** Four histograms display. Engine size clusters around 2.0–4.0 L, cylinders peak at 4, 6, and 8 combined fuel consumption is right-skewed, and CO<sub>2</sub> emissions form a roughly bell-shaped curve centered near 250 g/km.


Now I check the relationship between each feature and the target variable using scatter plots **to predict emissions.**

```
plt.scatter(cdf.FUELCONSUMPTION_COMB, cdf.CO2EMISSIONS, color='blue')
plt.xlabel("FUELCONSUMPTION_COMB")
plt.ylabel("Emission")
plt.show()
```



**Result:** A strong **positive linear relationship** - as fuel consumption increases, CO<sub>2</sub> emissions rise proportionally.

```
plt.scatter(cdf.ENGINESIZE, cdf.CO2EMISSIONS, color='blue')
plt.xlabel("Engine size")
plt.ylabel("Emission")
plt.show()
```




**Result:** Another positive trend, though with more scatter. **Larger engines generally produce more CO<sub>2</sub>**, but the **relationship is noisier** than fuel consumption.

## Exercise: Cylinders vs. Emission

Let's visualize how cylinder count correlates with CO<sub>2</sub> emissions.

```
plt.scatter(cdf.CYLINDERS, cdf.CO2EMISSIONS, color='blue')
plt.xlabel("Cylinders")
plt.ylabel("Emission")
plt.show()
```



**Result:** A clear upward trend - vehicles with more cylinders emit more CO<sub>2</sub>. The data clusters at specific cylinder counts (4, 6, 8, etc.), making this a reasonable but coarser predictor than engine size or fuel consumption.

---


## 4. Train/Test Split

I split the data so the model trains on 80% of observations and tests on the remaining 20%. This ensures I evaluate performance on data the model has never seen by creating a mask to select random rows using `np.random.rand()` function:

```
msk = np.random.rand(len(df)) < 0.8
train = cdf[msk]
test = cdf[~msk]
```

Let's visualize the training subset.

```
plt.scatter(train.ENGINESIZE, train.CO2EMISSIONS, color='blue')
plt.xlabel("Engine size")
plt.ylabel("Emission")
plt.show()
```



**Result:** The training data preserves the same positive trend observed in the full dataset.

---

## 5. Building the Simple Linear Regression Model

I use scikit-learn's `LinearRegression` to fit a line through the training data, minimizing the residual sum of squares.

```
from sklearn import linear_model

regr = linear_model.LinearRegression()
train_x = np.asanyarray(train[['ENGINE_SIZE']])
train_y = np.asanyarray(train[['CO2EMISSIONS']])
regr.fit(train_x, train_y)

print('Coefficients: ', regr.coef_)
print('Intercept: ', regr.intercept_)
```

```
Coefficients:  [[38.41986307]]
Intercept:    [127.24947105]
```

The fitted line is:

$$\text{CO}_2 \text{ Emissions} = 38.42 \times \text{Engine Size} + 127.25$$


Each additional litre of engine size adds roughly **39 g/km** of CO<sub>2</sub>, on average.

Plotting the regression line over the training data shows how well it captures the trend.

```
plt.scatter(train.ENGINE_SIZE, train.CO2EMISSIONS, color='blue')
```



```
plt.plot(train_x, regr.coef_[0][0]*train_x + regr.intercept_[0], '-r')
plt.xlabel("Engine size")
plt.ylabel("Emission")
plt.show()
```



**Result:** The red regression line cuts through the scatter, **rising steadily with engine size**. While some variance remains (points spread around the line), the overall linear trend is clear.

---

## 6. Evaluating the Model

I test the model on the held-out 20% and compute three standard metrics.

```
from sklearn.metrics import r2_score

test_x = np.asanyarray(test[['ENGINE SIZE']])
test_y = np.asanyarray(test[['CO2 EMISSIONS']])
test_y_ = regr.predict(test_x)

print("Mean absolute error: %.2f" % np.mean(np.absolute(test_y_ - test_y)))
print("Residual sum of squares (MSE): %.2f" % np.mean((test_y_ - test_y) ** 2))
print("R2-score: %.2f" % r2_score(test_y, test_y_))
```

Mean absolute error: 24.80  
 Residual sum of squares (MSE): 1034.49  
 R2-score: 0.77

- **MAE (24.80 g/km):** On average, the model's predictions are off by about 25 g/km.
- **MSE (1034.49):** Squared errors penalize larger deviations.



- **R<sup>2</sup> (0.77): Engine size** explains 77% of the variance in CO<sub>2</sub> emissions - a solid fit for a single feature. (best fit is **1.0**)
- 

## 7. Exercise: Using **Fuel Consumption** as a Feature

Now I repeat the entire workflow using **FUELCONSUMPTION\_COMB** as the predictor instead of engine size. This lets me compare which single feature performs better.

### Step 1 - Prepare the training and testing data:

```
train_x = train[["FUELCONSUMPTION_COMB"]]  
test_x = test[["FUELCONSUMPTION_COMB"]]
```

Lets see what the evaluation metrics are if we trained a regression model using the **FUELCONSUMPTION\_COMB** feature.

Start by selecting **FUELCONSUMPTION\_COMB** as the train\_x data from the **train** dataframe, then select **FUELCONSUMPTION\_COMB** as the test\_x data from the **test** dataframe

### Step 2 - Train the model:

```
regr = linear_model.LinearRegression()  
regr.fit(train_x, train_y)
```

▼ LinearRegression ⓘ ?

▼ Parameters

|                            |       |
|----------------------------|-------|
| <code>fit_intercept</code> | True  |
| <code>copy_X</code>        | True  |
| <code>tol</code>           | 1e-06 |
| <code>n_jobs</code>        | None  |
| <code>positive</code>      | False |

▼ Fitted attributes

| Name                           | Type                   | Value                    |
|--------------------------------|------------------------|--------------------------|
| <code>coef_</code>             | ndarray[float64](1, 1) | [[15.87]]                |
| <code>feature_names_in_</code> | ndarray[object](1,)    | ['FUELCONSUMPTION_COMB'] |
| <code>intercept_</code>        | ndarray[float64](1,)   | [71.82]                  |
| <code>n_features_in_</code>    | int                    | 1                        |
| <code>rank_</code>             | int                    | 1                        |
| <code>singular_</code>         | ndarray[float64](1,)   | [102.33]                 |

### Step 3 - Make predictions:

```
predictions = regr.predict(test_x)
```

### Step 4 - Evaluate:

```
print("Mean Absolute Error: %.2f" % np.mean(np.absolute(predictions - test_y)))
print("Mean Squared Error: %.2f" % np.mean((predictions - test_y) ** 2))
print("R2-score: %.2f" % r2_score(test_y, predictions))
```

Mean Absolute Error: 22.15

Mean Squared Error: 873.98

R2-score: 0.80

The MAE drops **from ~25 g/km (engine size) to ~22 g/km (fuel consumption)**, and  $R^2$  jumps from 0.77 to 0.80. Combined fuel consumption is a substantially stronger single predictor of CO<sub>2</sub> emissions - which makes intuitive sense, since fuel burned is directly tied to CO<sub>2</sub> produced.

---

That's the full walkthrough. I used scikit-learn to train and evaluate linear regression models predicting CO<sub>2</sub> emissions, and showed that `FUELCONSUMPTION_COMB` outperformed `ENGINE_SIZE` - **lowering error** and **raising  $R^2$**  (feature selection dramatically affects model performance()).